

# Project 3 - Classification Trees and Rules

---

Ola Tranum Arnegard - 12/08/2019

---

**Classification Trees or Rules Implementation:** Please implement the Classification Trees or Rules algorithm and submit the deliverables as specified in the project rubric document under Modules.

## Intro

Use Classification Trees (*c5.0*) to predict correctly from 3 flower **species** from 4 features given by **IRIS\_data**.

## Preprocessing

- Loaded data as **IRIS\_data** from same directory
- Randomized all **IRIS\_data** (rows)

This is what the data looks like:

```
head(IRIS_data,5)
```

```
##      sepal_length sepal_width petal_length petal_width      species
## 12             4.8         3.4         1.6         0.2      Iris-setosa
## 112            6.4         2.7         5.3         1.9  Iris-virginica
## 29             5.2         3.4         1.4         0.2      Iris-setosa
## 42             4.5         2.3         1.3         0.3      Iris-setosa
## 52             6.4         3.2         4.5         1.5  Iris-versicolor
```

## Implementation

### Create Model

```
train_IRIS <- as.data.frame(IRIS_data[1:125,-5]) #Train data (no names or labels)
train_labels <- as.data.frame(IRIS_data[1:125,5]) #Corresponding prediction labels

IRIS_model <- C5.0(train_IRIS, train_labels[,1]) #Create model
```

### Prediction

```
test_IRIS <- IRIS_data[126:nrow(IRIS_data),-5] #Corresponding to the training data
test_labels <- IRIS_data[126:nrow(IRIS_data),5] #Actual labels

IRIS_predictions <- predict(IRIS_model, test_IRIS) #Predicted labels

CrossTable(IRIS_predictions, test_labels,
            prop.chisq = FALSE, prop.t = FALSE, dnn = c('predicted', 'actual'))
```

```
##
##
##      Cell Contents
## |-----|
## |               N |
## |      N / Row Total |
## |      N / Col Total |
## |-----|
##
##
## Total Observations in Table:  25
##
##
##      predicted | actual
##      predicted |      Iris-setosa | Iris-versicolor | Iris-virginica |      Row Total |
## -----|-----|-----|-----|-----|
##      Iris-setosa |      8 |      0 |      0 |      8 |
##                  |      1.000 |      0.000 |      0.000 |      0.320 |
##                  |      0.800 |      0.000 |      0.000 |      |
## -----|-----|-----|-----|-----|
##      Iris-versicolor |      2 |      5 |      0 |      7 |
##                  |      0.286 |      0.714 |      0.000 |      0.280 |
##                  |      0.200 |      0.625 |      0.000 |      |
## -----|-----|-----|-----|-----|
##      Iris-virginica |      0 |      3 |      7 |      10 |
##                  |      0.000 |      0.300 |      0.700 |      0.400 |
##                  |      0.000 |      0.375 |      1.000 |      |
## -----|-----|-----|-----|-----|
##      Column Total |      10 |      8 |      7 |      25 |
##                  |      0.400 |      0.320 |      0.280 |      |
## -----|-----|-----|-----|-----|
##
##
```

## Results

We have 3 different flowers to predict. As we see in the crosstable above, 2 *Iris-versicolor* and 3 *Iris-virginica* are predicted wrong, with an accuracy of 71% and 70% respectively: 2 predicted *Iris-versicolor* should have been 2 *Iris-setosa* and 3 *Iris-virginica* should have been *Iris-versicolor*.

### Decision tree

- From start to completion, our model will continuously calculate the split that has the highest *information gain* and use that as the next split, i.e. the split that gives the highest difference in **entropy**.

$$inf.gain = \sum_{before\ split} -p_i \log_2 p_i - \sum_{after\ split} -p_i \log_2 p_i$$

- It is important to not split excessively because it may lead to overfitting. However, as for most classification algorithms and predictors, it's always a tradeoff between having a model that is overly generalized and a model that is overfitted. To solve this, we often use either *pre-pruning* or *post-pruning*, that is, respectively, give a fixed number of allowed splits before termination or complete the algorithm and afterwards remove insignificant splits.
- Our model uses *post-pruning*
- Below we can see a summary of our model: only 3 splits long but with a accuracy of 98.4% on our training data. This is substantially higher than our predictions, and therefore, one may wonder if the model is overfitted. However, since we only use 3 splits, it is most likely not.
- Most of the wrong predictions come from *Iris-virginica* that really are *Iris-versicolor*. Not surprisingly, we can see from our model that all of the errors come from the same mistake.

```
summary(IRIS_model)
```

```
##
## Call:
## C5.0.default(x = train_IRIS, y = train_labels[, 1])
##
##
## C5.0 [Release 2.07 GPL Edition]      Sat Oct 12 15:02:57 2019
## -----
##
## Class specified by attribute `outcome'
##
## Read 125 cases (5 attributes) from undefined.data
##
## Decision tree:
##
## petal_width <= 0.4: Iris-setosa (40)
## petal_width > 0.4:
##   ...petal_length > 4.8: Iris-virginica (41/1)
##     petal_length <= 4.8:
##       ...petal_width <= 1.6: Iris-versicolor (40)
##         petal_width > 1.6: Iris-virginica (4/1)
##
##
## Evaluation on training data (125 cases):
##
##      Decision Tree
##      -----
##      Size      Errors
##
##      4      2( 1.6%)  <<
##
##
##      (a)  (b)  (c)    <-classified as
##      ----  ----  ----
```

```
##      40          (a): class Iris-setosa
##      40      2    (b): class Iris-versicolor
##      43      (c): class Iris-virginica
##
##
## Attribute usage:
##
## 100.00% petal_width
##  68.00% petal_length
##
##
## Time: 0.0 secs
```